

ECE253, HW3, Report

William Argus A12802324

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind. By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

Problem 1

In []:

Matlab Code

```
%% ECE253, HW3, Problem 1, Version 2
% William Argus A12802324
%% setup
clc;
clear all;
close all;

image = imread('geisel.jpg');
imageGray = rgb2gray(image);
te = 130
edges = cannyEdgeDetection(imageGray, te);
figure(3);
imshow(im2uint8(edges/255));
title('Final edge image after thresholding', 'FontSize', 16);
%% function
function imageEdges = cannyEdgeDetection(imageGray, te)
    %% Part 1
    k = (1/159)*[2 4 5 4 2; 4 9 12 9 4; 5 12 15 12 5; 4 9 12 9 4
; 2 4 5 4 2];
    imageSmooth = conv2(imageGray, k);
    imageSmooth = im2uint8(imageSmooth/255);

    %% Part 2
```

%% Part 2

```
kx = [-1 0 1; -2 0 2; -1 0 1];  
ky = [-1 -2 -1; 0 0 0; 1 2 1];
```

```
gX = conv2(imageSmooth, kx);  
gY = conv2(imageSmooth, ky);  
absGradient = sqrt(gX.^2 + gY.^2);  
gX(gX == 0) = 0.0001;  
angleGradient = atan2(gY./gX);  
figure(1);  
imshow(im2uint8(absGradient/255));  
title('Original gradient magnitude image', 'FontSize', 16);
```

%% Part 3

```
angleGradient(-22.5 <= angleGradient & 22.5 > angleGradient)  
= 0;  
angleGradient(-67.5 <= angleGradient & -22.5 > angleGradient  
) = -45;  
angleGradient(67.5 > angleGradient & 22.5 <= angleGradient)  
= 45;  
angleGradient(-67.5 > angleGradient) = 90;  
angleGradient(67.5 <= angleGradient) = 90;
```

%for each pixel

```
sz = size(absGradient);
```

```
pixel1=0;
```

```
pixel2=0;
```

```
for i = 2:sz(1)-1
```

```
    for j = 2:sz(2)-1
```

```
        %get gradient direction to pick comparison pixels
```

```
        if angleGradient(i,j) == 90
```

```
            pixel1 = absGradient(i-1,j);
```

```
            pixel2 = absGradient(i+1,j);
```

```
        elseif angleGradient(i,j) == -45
```

```
            pixel1 = absGradient(i-1,j-1);
```

```
            pixel2 = absGradient(i+1,j+1);
```

```
        elseif angleGradient(i,j) == 0
```

```
            pixel1 = absGradient(i,j+1);
```

```
            pixel2 = absGradient(i,j-1);
```

```
        elseif angleGradient(i,j) == 45
```

```
            pixel1 = absGradient(i-1,j+1);
```

```
            pixel2 = absGradient(i+1,j-1);
```

```
        else
```

```
            print("ERROR!!");
```

```

        end
        if absGradient(i,j) > pixel1 && absGradient(i,j) > pixel2
            nothing=0;
        else
            absGradient(i,j)=0;
        end
    end
end

figure(2);
imshow(im2uint8(absGradient/255));
title('Image after NMS', 'FontSize', 16);

%% Part 4

absGradient(absGradient < te) =0;
imageEdges = absGradient;
end

```

In [194]:

```

#Results
import matplotlib.pyplot as plt

P1_1 = plt.imread('P1.1.png')
P1_2 = plt.imread('P1.2.png')
P1_3 = plt.imread('P1.3.png')

plt.figure(figsize = (10,10))
plt.imshow(P1_1)
plt.figure(figsize = (10,10))
plt.imshow(P1_2)
plt.figure(figsize = (10,10))
plt.imshow(P1_3)

print('Value of te used to product final image: 130')
print('This ')

```

Value of te used to product final image: 130

Original gradient magnitude image

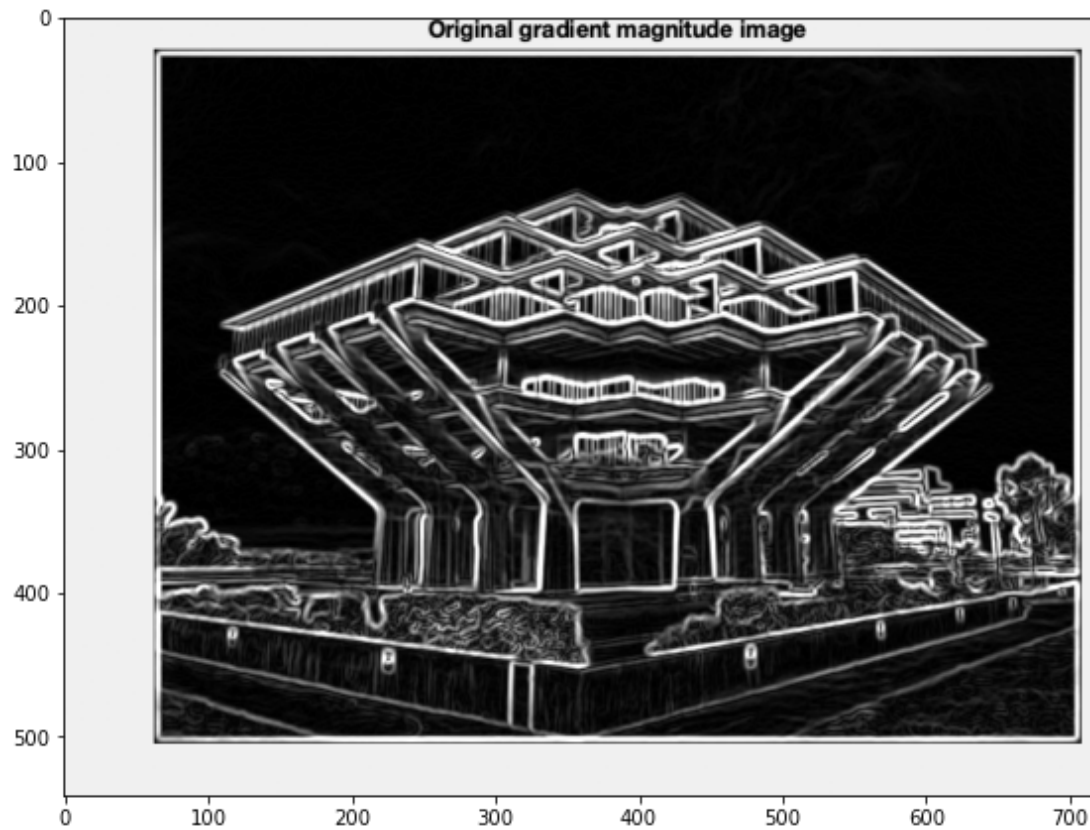

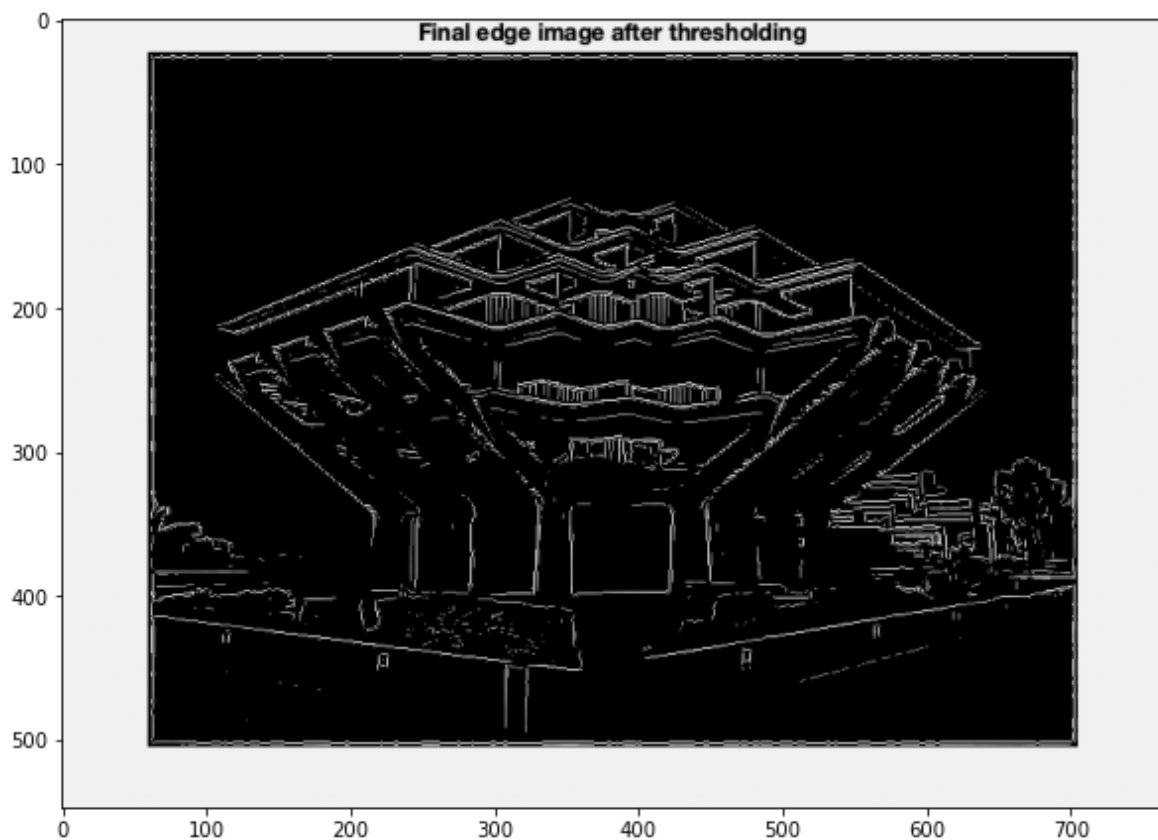


Image after NMS





Problem 2

Problem 2 (i)

In []:

```
### Matlab Code for (i)

%% ECE253, HW3, Problem 2i, Version 1
% William Argus A12802324
%% setup
clc;

image = imread('Car.tif');
figure(1);
imshow(image);
title('Unpadded original image')
colorbar;
%%
sz = size(image);
padsizel = (512-sz(1))/2 +1;
padsizel2 = (512-sz(2))/2 +1;
```

```

paddedImage = im2uint8(zeros(512,512));

paddedImage((padsizel+1):(padsizel+sz(1)), (padsizel2+1):(padsizel2+sz(2))) = image;
%figure(2);
%imshow(paddedImage);

imFFT = fft2(paddedImage);
%test = ifft2(imFFT);
imFFT = fftshift(imFFT);
%figure(4);
%imshow(im2uint8(test/255));
figure(2)
imagesc(-256:255,-256:255,log(abs(imFFT)));
colorbar;
title('2D DFT log-magnitude of original image');
xlabel('u');
ylabel('v');

[u,v] = meshgrid(-256:255);

%iterate through all possible u values, then through all possible v values,
%then for each (u,v), calculate Dk
%then Hnr for each (u,v)?
%since there is 4 bursts, need 4 pairs of coordinates for the bursts

%calculate Hnr
Hnr = ones(512,512);
Do = 20
n = 2
uv = [83, -175; 83, -90; 90, 81; 90, 170]
for k = 1:4
    uk = uv(k,1);
    vk = uv(k,2);
    %Dpk = zeros(512,512);
    %Dmk = zeros(512,512);
    Dpk = ((u - uk).^2 + (v - vk).^2).^0.5;
    Dmk = ((u + uk).^2 + (v + vk).^2).^0.5;

    term1 = ( 1+ (Do.*(Dpk.^-1)).^(2*n) ).^-1;
    term2 = ( 1+ (Do.*(Dmk.^-1)).^(2*n) ).^-1;
    temp = term1.*term2;

```

```

    Hnr = Hnr.*temp;
end

figure(3);
imshow(Hnr);
colorbar;
title('The butterworth Notch Reject Filter in frequency domain H
nr(u,v)');
h = gca;
set(h, 'Visible', 'On')

freqImage = Hnr.*imFFT;
figure(4);
imagesc(-256:255,-256:255,log(abs(freqImage)));
colorbar;
title('The frequency domain of the image after filtering');
xlabel('u');
ylabel('v');

unshift = ifftshift(freqImage);
resultPadded = ifft2(unshift);
%unpad image
result = resultPadded((padsizel+1):(padsizel+sz(1)), (padsizel+1
):(padsizel+sz(2)));
figure(5);
imshow(im2uint8(result/255));
colorbar;
title('The final filtered image');

```

In [206]:

```
#Results
```

```
import matplotlib.pyplot as plt
```

```
P2i_1 = plt.imread('P2i.1.png')
```

```
P2i_2 = plt.imread('P2i.2.png')
```

```
P2i_3 = plt.imread('P2i.3.png')
```

```
P2i_4 = plt.imread('P2i.4.png')
```

```
P2i_5 = plt.imread('P2i.5.png')
```

```
plt.figure(figsize = (20,20))
```

```
plt.subplot(321)
```

```
plt.imshow(P2i_1)
```

```
plt.title('Unpadded original image')
```

```
plt.subplot(322)
```

```
plt.imshow(P2i_2)
```

```
plt.title('The corresponding 2D DFT log-magnitude')
```

```
plt.subplot(323)
```

```
plt.imshow(P2i_3)
```

```
plt.title('The butterworth Notch Reject Filter in frequency domain  
in HNR(u,v)')
```

```
plt.subplot(324)
```

```
plt.imshow(P2i_4)
```

```
plt.title('The frequency domain of image after filtering')
```

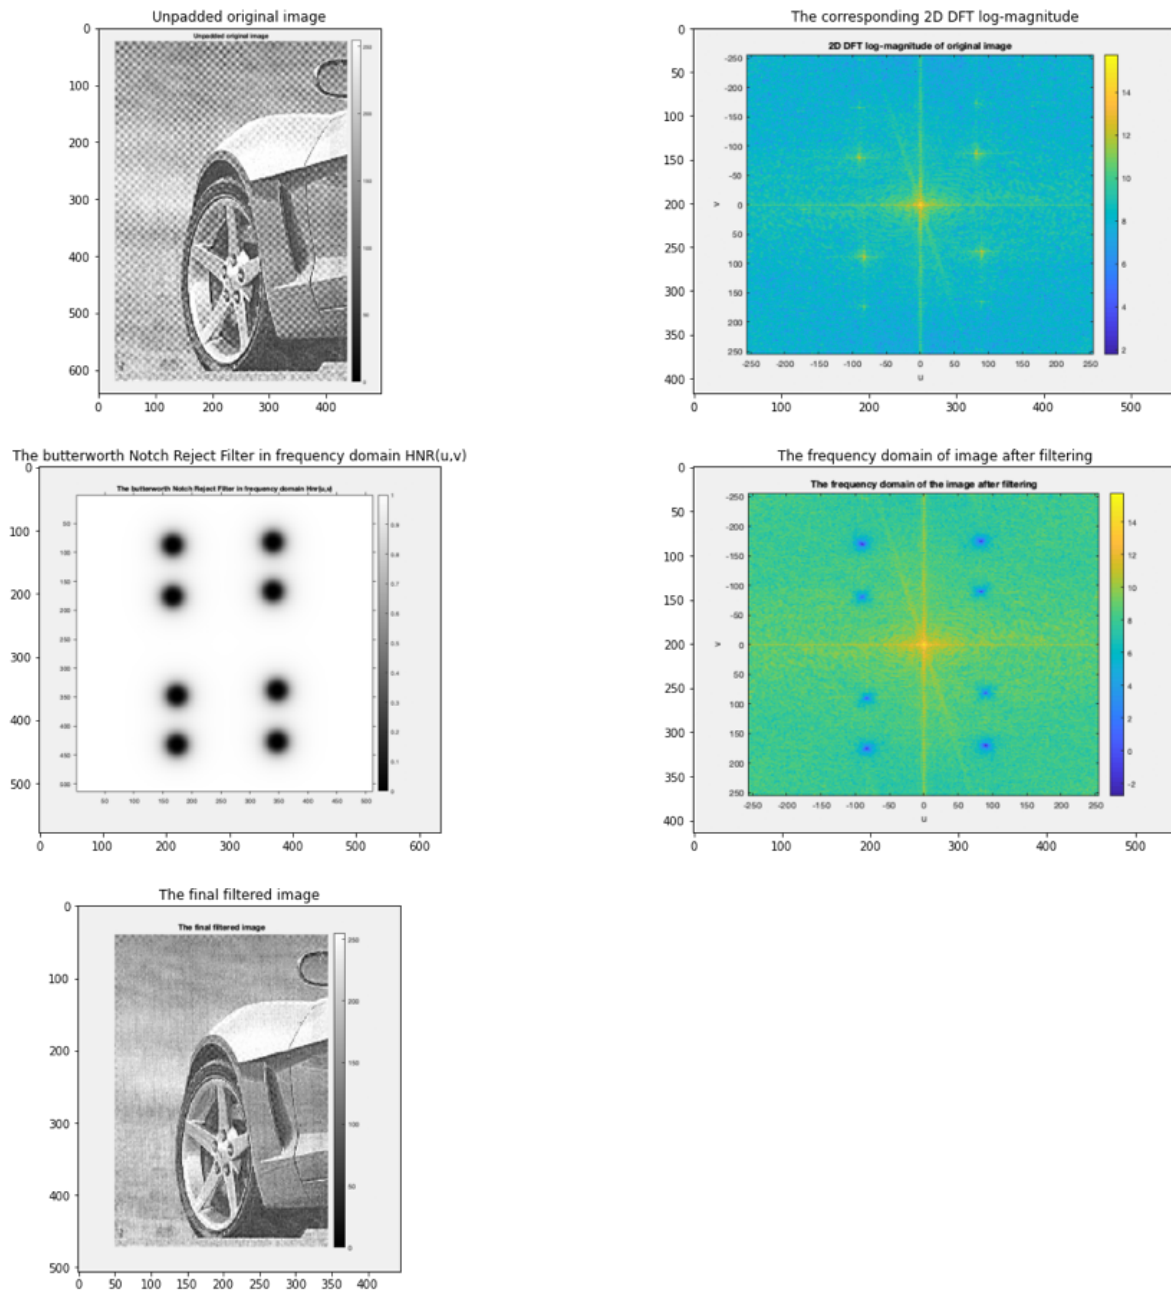
```
plt.subplot(325)
```

```
plt.imshow(P2i_5)
```

```
plt.title('The final filtered image')
```


Out[206]:

Text(0.5, 1.0, 'The final filtered image')



Do = 20 n = 2 u, v = (8 terms below) u1 = 83, v1 = -175 u2 = 83, v2 = -90 u3 = 90, v3 = 81
u4 = 90, v4 = 170

Problem 2 (ii)

In []:

```
### Matlab code for (ii)
```

```
%% ECE253, HW3, Problem 2ii, Version 1  
% William Argus A12802324
```

```

%% setup

image = imread('Street.png');
figure(1);
imshow(image);
title('Unpadded original image')
colorbar;

sz = size(image);
padsizel = (512-sz(1))/2 +1;
padsizel2 = (512-sz(2))/2 +1;
paddedImage = im2uint8(zeros(512,512));

paddedImage((padsizel+1):(padsizel+sz(1)), (padsizel2+1):(padsizel2+sz(2))) = image;
%figure(2);
%imshow(paddedImage);

imFFT = fft2(paddedImage);
%test = ifft2(imFFT);
imFFT = fftshift(imFFT);
%figure(4);
%imshow(im2uint8(test/255));
figure(2)
imagesc(-256:255,-256:255,log(abs(imFFT)));
colorbar;
title('2D DFT log-magnitude of original image');
xlabel('u');
ylabel('v');

[u,v] = meshgrid(-256:255);
%%
%iterate through all possible u values, then through all possible v values,
%then for each (u,v), calculate Dk
%then Hnr for each (u,v)?
%since there is 4 bursts, need 4 pairs of coordinates for the bursts

%calculate Hnr
Hnr = ones(512,512);
Do = 50
n = 2
uv = [0, -165; 165, 0;]

```

```

for k = 1:2
    uk = uv(k,1);
    vk = uv(k,2);
    Dpk = ((u - uk).^2 + (v - vk).^2).^5;
    Dmk = ((u + uk).^2 + (v + vk).^2).^5;

    term1 = ( 1+ (Do.*(Dpk.^-1)).^(2*n) ).^-1;
    term2 = ( 1+ (Do.*(Dmk.^-1)).^(2*n) ).^-1;
    temp = term1.*term2;
    Hnr = Hnr.*temp;
end

figure(3);
imshow(Hnr);
colorbar;
title('The butterworth Notch Reject Filter in frequency domain H
nr(u,v)');
h = gca;
set(h, 'Visible', 'On')

freqImage = Hnr.*imFFT;
figure(4);
imagesc(-256:255,-256:255,log(abs(freqImage)));
colorbar;
title('The frequency domain of the image after filtering');
xlabel('u');
ylabel('v');

unshift = ifftshift(freqImage);
resultPadded = ifft2(unshift);
%unpad image
result = resultPadded((padsizel+1):(padsizel+sz(1)), (padsizel+1
):(padsizel+sz(2)));
figure(5);
imshow(im2uint8(result/255));
colorbar;
title('The final filtered image');

```

In [207]:

```
#Results
```

```
import matplotlib.pyplot as plt
```

```
P2ii_1 = plt.imread('P2ii.1.png')
```

```
P2ii_2 = plt.imread('P2ii.2.png')
```

```
P2ii_3 = plt.imread('P2ii.3.png')
```

```
P2ii_4 = plt.imread('P2ii.4.png')
```

```
P2ii_5 = plt.imread('P2ii.5.png')
```

```
plt.figure(figsize = (20,20))
```

```
plt.subplot(321)
```

```
plt.imshow(P2ii_1)
```

```
plt.title('Unpadded original image')
```

```
plt.subplot(322)
```

```
plt.imshow(P2ii_2)
```

```
plt.title('The corresponding 2D DFT log-magnitude')
```

```
plt.subplot(323)
```

```
plt.imshow(P2ii_3)
```

```
plt.title('The butterworth Notch Reject Filter in frequency domain  
in HNR(u,v)')
```

```
plt.subplot(324)
```

```
plt.imshow(P2ii_4)
```

```
plt.title('The frequency domain of image after filtering')
```

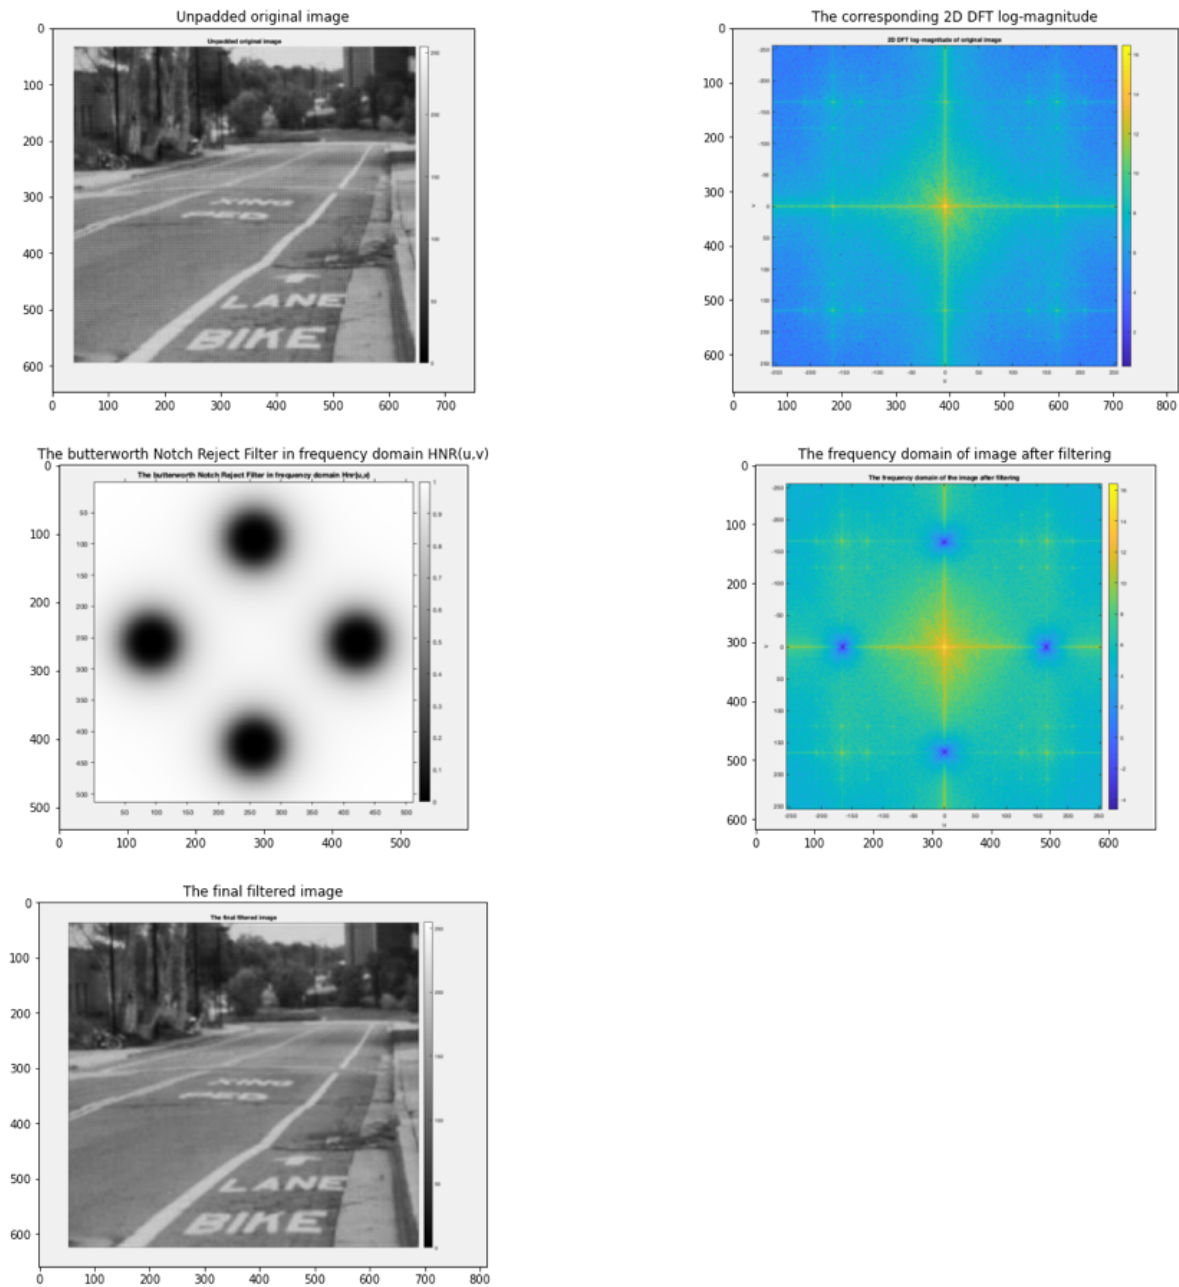
```
plt.subplot(325)
```

```
plt.imshow(P2ii_5)
```

```
plt.title('The final filtered image')
```

Out[207]:

```
Text(0.5, 1.0, 'The final filtered image')
```



$D_0 = 50$ $n = 2$ $uv = (4 \text{ terms below})$ $u_1 = 0$, $v_1 = -165$ $u_2 = 165$, $v_2 = 0$

In []:

Problem 3

In [16]:

```
import torch
import torchvision
import torchvision.transforms as transforms
```

In [17]:

```
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((.5, .5, .5),
                                                       (.5, .5, .5))])
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True,
                                         transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

Files already downloaded and verified


In [18]:

```
import matplotlib.pyplot as plt
import numpy as np
# functions to show an image

def imshow(img):
    img = img / 2 + 0.5      # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.next()

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```



plane cat deer car

In [19]:

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
```

In [20]:

```
import torch.optim as optim

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```


In [23]:

```
for epoch in range(2): # loop over the dataset multiple times
    print('Training, epoch: ', epoch)
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running_loss = 0.0

print('Finished Training')
```

```
Training, epoch:  0
[1,  2000] loss: 2.179
[1,  4000] loss: 1.909
[1,  6000] loss: 1.702
[1,  8000] loss: 1.596
[1, 10000] loss: 1.541
[1, 12000] loss: 1.488
Training, epoch:  1
[2,  2000] loss: 1.400
[2,  4000] loss: 1.408
[2,  6000] loss: 1.368
[2,  8000] loss: 1.338
[2, 10000] loss: 1.314
[2, 12000] loss: 1.287
Finished Training
```

In [24]:

```
PATH = './cifar_net.pth'  
torch.save(net.state_dict(), PATH)
```

In [25]:

```
dataiter = iter(testloader)  
images, labels = dataiter.next()  
  
# print images  
imshow(torchvision.utils.make_grid(images))  
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j  
in range(4)))
```



GroundTruth: cat ship ship plane

In [26]:

```
net = Net()  
net.load_state_dict(torch.load(PATH))
```

Out[26]:

<All keys matched successfully>

In [27]:

```
outputs = net(images)
```

In [29]:

```
_, predicted = torch.max(outputs, 1)
print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
                                for j in range(4)))
```

Predicted: bird car car plane

In [30]:

```
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy of the network on the 10000 test images: %d %%'
      % (
          100 * correct / total))
```

Accuracy of the network on the 10000 test images: 53
%

In [31]:

```
class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))
```

```
Accuracy of plane : 73 %
Accuracy of   car : 68 %
Accuracy of  bird : 38 %
Accuracy of   cat : 45 %
Accuracy of  deer : 49 %
Accuracy of   dog : 25 %
Accuracy of  frog : 60 %
Accuracy of horse : 66 %
Accuracy of  ship : 53 %
Accuracy of truck : 54 %
```

In [32]:

```
device = torch.device("cuda:0" if torch.cuda.is_available() else
    "cpu")

# Assuming that we are on a CUDA machine, this should print a CU
DA device:

print(device)
```

cuda:0

In [33]:

```
net.to(device)
```

Out[33]:

```
Net(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(
1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding
=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=
(1, 1))
  (fc1): Linear(in_features=400, out_features=120, b
ias=True)
  (fc2): Linear(in_features=120, out_features=84, bi
as=True)
  (fc3): Linear(in_features=84, out_features=10, bia
s=True)
)
```

In [34]:

```
inputs, labels = data[0].to(device), data[1].to(device)
```

In []:


In []:



In []:


In [67]:

```
numTraining = 0
inputSize = 0
inputsCheck = 0
for i, data in enumerate(trainloader, 0):
    # get the inputs; data is a list of [inputs, labels]
    inputs, labels = data
    inputsCheck = inputs
    inputSize = len(inputs)
    numTraining = numTraining + 1
print('Number of minibatches: ', numTraining)
print('Number of images per batch: ',inputSize)
imshow(inputsCheck[0])
imshow(inputsCheck[1])
imshow(inputsCheck[2])
imshow(inputsCheck[3])
```

Number of minibatches: 12500
Number of images per batch: 4







(ii)

How many images and batches are used to train the network?

12,500 mini-batches loaded in with trainloader. 4 training images per mini-batch, as found in the cell above. Therefore, a total of 50,000 images used to train the network. 1 batch in total since it was mentioned in lecture that the batch refers to all of the images in the training dataset, therefore in this case the entirety of the training set is one batch divided into mini batches with 4 images in each to be used to train the network.

(iii)

Do we normalize the images? What do we do in the example?

Yes, normalization of the training and test set is necessary in a neural network as the ranges of feature distributions must be the same in order for the network to learn the correct weights. This is because the weights are determined through multiplication of the errors found via backpropagation by the defined learning rate of the network. In the above example, the images are indeed normalized, specifically in the definition of 'transform' (copied below), which is used as an input parameter when loading the training and testing datasets from CIFAR10.

```
"transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((.5, .5, .5),(.5, .5, .5))])"
```

(iv)

The losses are dropping! Can you plot out the training loss?

Yes, see plots below for epoch 1 and 2. Please note that training took hours so rather than rewriting the training code block to save the losses in a vector or list for plotting later, the vector was defined by hand from the printed out losses in the training block in order to save time and sanity. It can be observed that the loss decreases rapidly before starting to level off towards the end of the second epoch. It briefly increases at the beginning of the second epoch before decreasing again. This is not uncommon in training.

In [65]:

```
import numpy as np
import matplotlib.pyplot as plt

#epoch 1
losses1 = np.array([2.179,1.909,1.702,1.596,1.541,1.488])
#epoch 2
losses2 = np.array([1.400,1.408,1.368,1.338,1.314,1.287])
#mini-batch number
miniBatch = np.array([2000,4000,6000,8000,10000,12000])

fig=plt.figure(1)
plt.scatter(miniBatch, losses1)
plt.title('Losses vs mini batch number, epoch 1')
plt.xlabel('mini batch number')
plt.ylabel('loss')
plt.ylim([1.25,2.2])


fig=plt.figure(2)
plt.scatter(miniBatch, losses2)
plt.title('Losses vs mini batch number, epoch 2')
plt.xlabel('mini batch number')
plt.ylabel('loss')
plt.ylim([1.25,2.2])

fig=plt.figure(3)
lossesTotal = np.append(losses1, losses2)

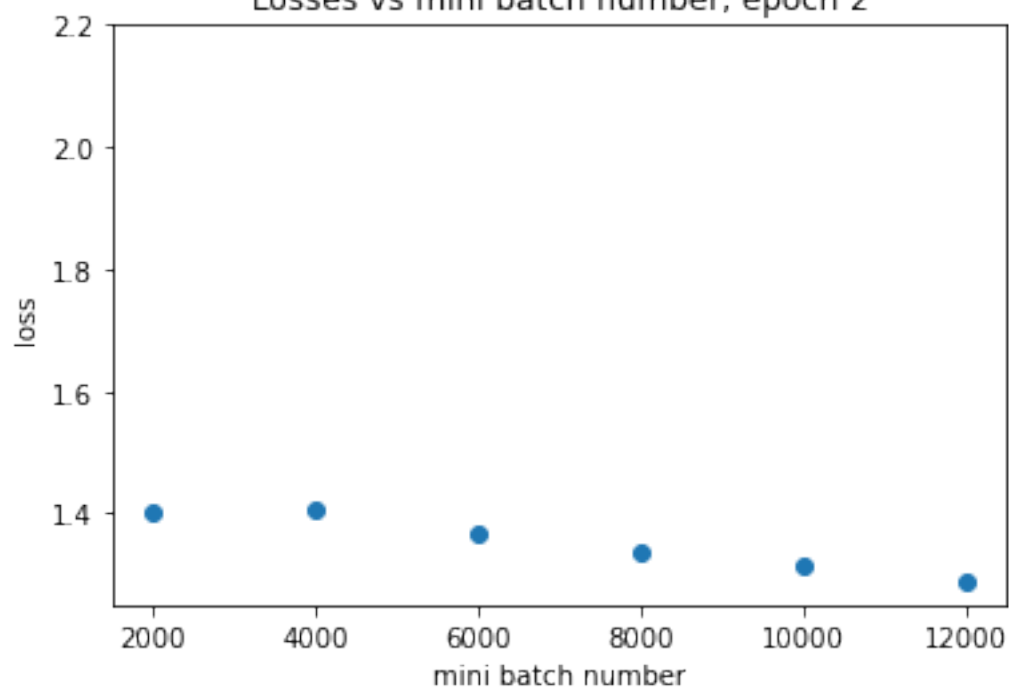
miniBatchTotal = np.append(miniBatch, miniBatch+12500)
plt.scatter(miniBatchTotal, lossesTotal)
plt.title('Losses vs number of mini batches, epochs 1 and 2')
plt.xlabel('number of mini batches')
plt.ylabel('loss')
plt.ylim([1.25,2.2])

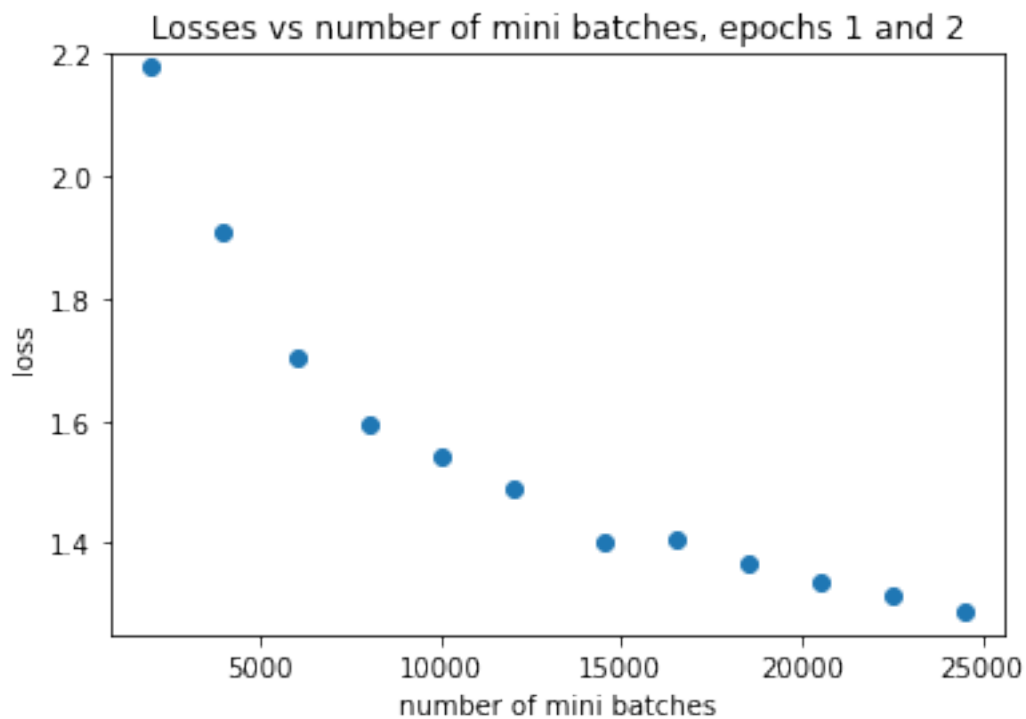
plt.show()
```

Losses vs mini batch number, epoch 1



Losses vs mini batch number, epoch 2





In [64]:

```
print(miniBatch+12500)
```

```
[14500 16500 18500 20500 22500 24500]
```

(v)

Now the network is done training. Can you check some successful cases and some failure cases (show some images classified by the network)?

Yes, see below, using code from the tutorial above. The network is not perfect but is much better than a random guess, showing that the network did learn something.

In [88]:

```
net = Net()
net.load_state_dict(torch.load(PATH))

dataiter = iter(testloader)
for i in range(0,10):
    images, labels = dataiter.next()

    # print images
    imshow(torchvision.utils.make_grid(images))
    print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] f
or j in range(4)))


    outputs = net(images)

    _, predicted = torch.max(outputs, 1)

    print('Predicted: ', ' '.join('%5s' % classes[predicted[j]]
for j in range(4)))
```



GroundTruth: cat ship ship plane
Predicted: bird car car plane



GroundTruth: frog frog car frog
Predicted: cat frog cat deer



GroundTruth: cat car plane truck

Predicted: cat car plane truck



GroundTruth: dog horse truck ship

Predicted: cat horse truck frog




GroundTruth: dog horse ship frog

Predicted: cat cat ship frog




GroundTruth: horse plane deer truck

Predicted: horse plane plane truck



GroundTruth: dog bird deer plane

Predicted: deer frog deer bird



GroundTruth: truck frog frog dog

Predicted: car frog cat bird



GroundTruth: deer dog truck bird

Predicted: deer cat truck cat



GroundTruth: deer car truck dog

Predicted: horse car truck dog


(vi)

Can you visualize the output of the 1st layer of CNN using one image from the training set?

In [183]:

```
for i, data in enumerate(trainloader, 0):  
    # get the inputs; data is a list of [inputs, labels]  
    inputs, labels = data  
    image = inputs[0]  
    label = labels[0]  
print('Image to be used to visualize output of 1st layer of CNN'  
)  
imshow(image)  
print(label)
```

Image to be used to visualize output of 1st layer of
CNN



tensor(6)

In [211]:

```
from scipy.signal import convolve
from scipy.ndimage import correlate
def convolve3D (image, kernal):
    #inputs both (x,x,3)
    layer1 = correlate(image[:,:,:0], kernal[:,:,:0], mode = 'wrap
')
    layer2 = correlate(image[:,:,:1], kernal[:,:,:1], mode = 'wrap
')
    layer3 = correlate(image[:,:,:2], kernal[:,:,:2], mode = 'wrap
')
    result = np.dstack([layer1, layer2, layer3])
    return result
```

In [212]:

```
weights = net.conv1.weight.data.numpy()
print(weights.shape)
weights.T.shape
print(weights.shape)
newWeights = (1/2) + weights/2

imageNP = np.transpose(image.numpy(),(1,2,0))
imageNP = (1/2) + imageNP/2

output = [0,0,0,0,0,0]

for i in range(0,6):
    filter = np.transpose(newWeights[i,...],(1,2,0))
    output[i] = convolve3D(imageNP, filter)
    output[i] = ((output[i]-np.amin(output[i]))/(np.amax(output[
i])-np.amin(output[i])))

plt.figure(figsize=(10,10))
plt.subplot(321)
plt.imshow(output[0])
plt.title('output ' + str(1) + ' for first layer of CNN')

plt.subplot(322)
plt.imshow(output[1])
plt.title('output ' + str(2) + ' for first layer of CNN')

plt.subplot(323)
```

```
plt.imshow(output[2])

plt.title('output ' + str(3) + ' for first layer of CNN')

plt.subplot(324)
plt.imshow(output[3])
plt.title('output ' + str(4) + ' for first layer of CNN')

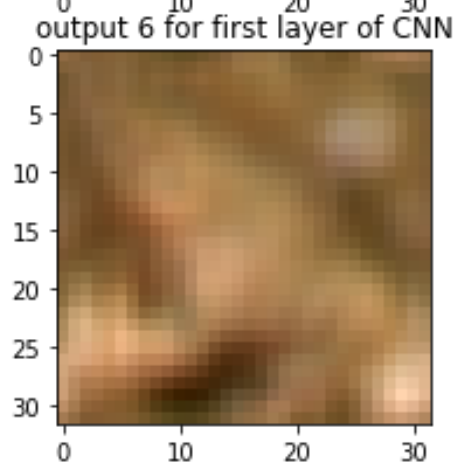
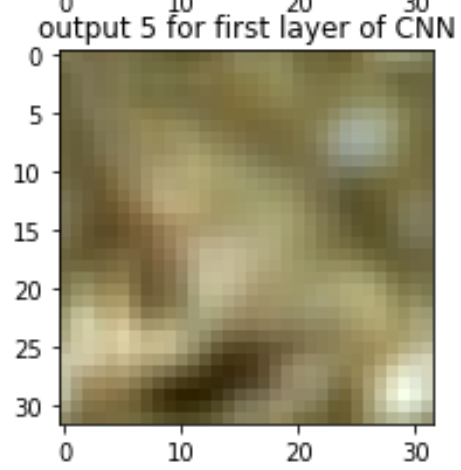
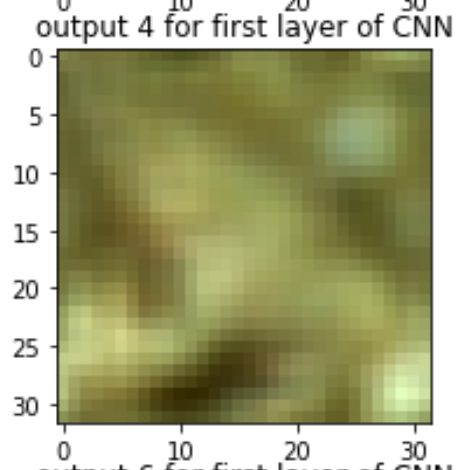



plt.subplot(325)
plt.imshow(output[4])
plt.title('output ' + str(5) + ' for first layer of CNN')

plt.subplot(326)
plt.imshow(output[5])
plt.title('output ' + str(6) + ' for first layer of CNN')
```

```
(6, 3, 5, 5)
(6, 3, 5, 5)
```

Out[212]:

```
Text(0.5, 1.0, 'output 6 for first layer of CNN')
```



In []: